

Spatial Filtering

Srikumar Ramalingam
School of Computing
University of Utah

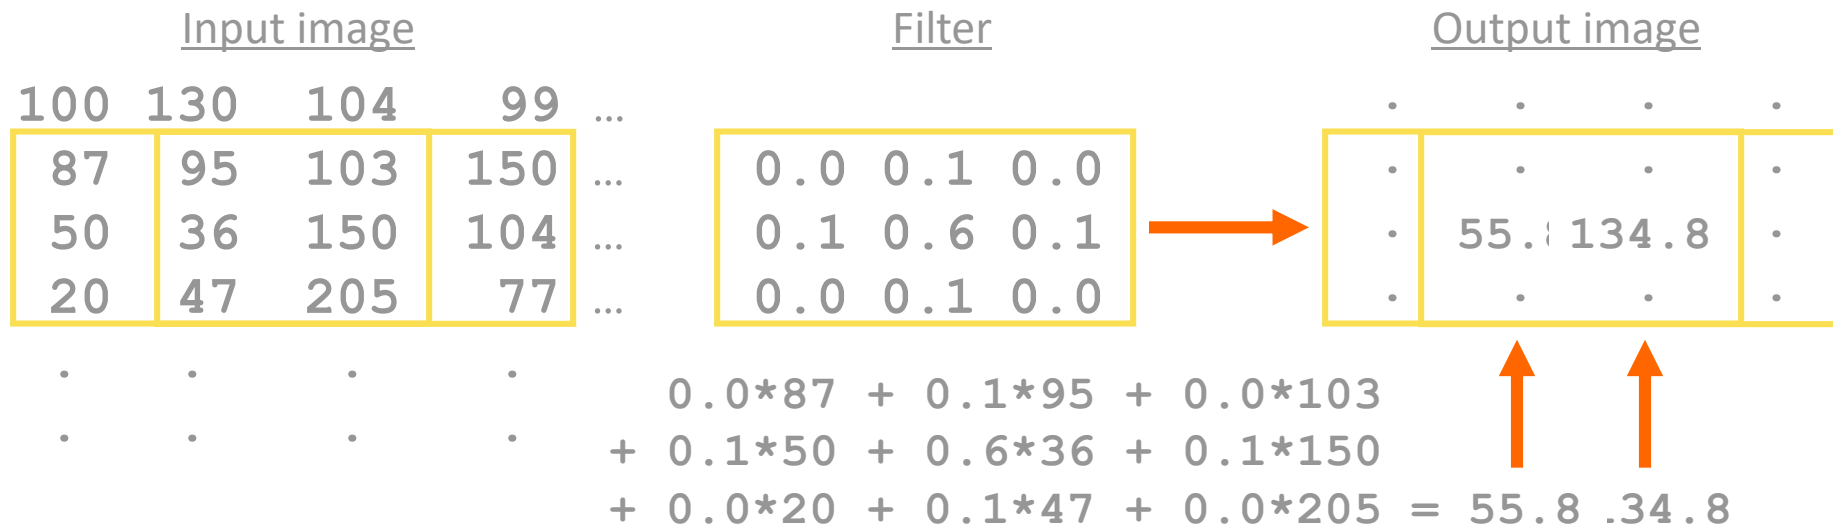
[Slides borrowed from Ross Whitaker]

Overview

- Correlation and convolution
- Linear filtering
 - Smoothing, kernels, models
 - Detection
 - Derivatives
- Nonlinear filtering
 - Median filtering
 - Bilateral filtering
 - Neighborhood statistics and nonlocal filtering

Cross Correlation

- Operation on image neighborhood and small ...
 - “mask”, “filter”, “stencil”, “kernel”
- Linear operations within a moving window




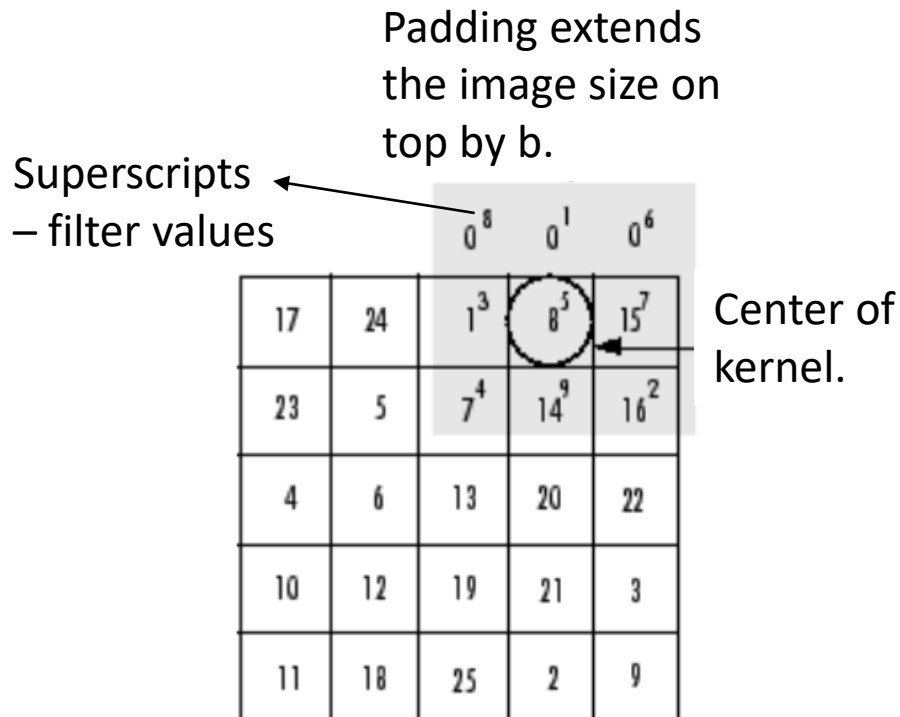
Cross Correlation

- 1D $g(x) = \sum_{s=-a}^a w(s) f(x + s)$
- 2D $g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$

$$w(s, t) = \begin{array}{ccccccc} & w(-a, -b) & \cdots & & \cdots & & w(a, -b) \\ & \vdots & & & & & \vdots \\ & & & \cdots & w(0, 0) & \cdots & \\ & \vdots & & & & & \vdots \\ & w(-a, b) & \cdots & & \cdots & & w(a, b) \end{array}$$

Correlation: Technical Details

- Boundary conditions – Pad image with amount (a,b) for a filter of size $(2a+1, 2b+1)$.




Correlation: Technical Details


- Boundary conditions – zero padding, replication of boundary pixels

Outside pixels are assumed to be 0's.

Superscripts – filter values





The pixel values are replicated from boundary pixels.



Correlation: Technical Details

- Boundary conditions
 - Pad image with amount (a,b)
 - Constant value or repeat edge values
 - Cyclical boundary conditions
 - Wrap or mirroring



Correlation: Technical Details

- **Boundaries**
 - Can also modify kernel – no long correlation
- **For analysis**
 - Image domains infinite
 - Data compact (goes to zero far away from origin)

$$g(x, y) = \sum_{s=-\infty}^{\infty} \sum_{t=-\infty}^{\infty} w(s, t) f(x + s, y + t)$$

Correlation: Properties

- Shift invariant – discrete equivalent of time invariant system - If $x(t)$ gives $y(t)$, then $x(t + t_0)$ gives $y(t + t_0)$.

$$g = w \circ f \quad \text{New notation for correlation}$$

$$w \circ f(x-x_0, y-y_0) = \sum_{s=-\infty}^{\infty} \sum_{t=-\infty}^{\infty} w(s, t) f(x-x_0+s, y-y_0+t) = g(x-x_0, y-y_0)$$

- Linear $w \circ (\alpha e + \beta f) = \alpha w \circ e + \beta w \circ f$

$$C_{wf} = w \circ f \quad \text{Compact notation}$$

Cross Correlation Continuous Case

- f, w must be “integrable”
 - Must die off fast enough so that integral is finite

$$g(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} w(s, t) f(x + s, y + t) ds dt$$

- Same properties as discrete case
 - Linear
 - Shift invariant

Filters: Considerations

- Normalize
 - Sums to one
 - Sums to zero (some cases, later)
- Symmetry
 - Left, right, up, down
 - Rotational
- Special case: auto correlation

$$C_{ff} = f \circ f$$

Examples 1



0	0	0
0	1	0
0	0	0



	1	1	1
$1/9 *$	1	1	1
	1	1	1



Examples 2



$$\frac{1}{9} * \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



$$\frac{1}{25} * \begin{matrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{matrix}$$



Smoothing and Noise

Noisy image




5x5 box filter




Other Filters

- Disk
 - Circularly symmetric, jagged in discrete case



- Gaussians
 - Circularly symmetric, smooth for large enough std-dev.
 - Must normalize in order to sum to one

$$\frac{1}{2\pi\sigma^2} \exp -\frac{(x^2 + y^2)}{2\sigma^2}$$



- Derivatives – discrete/finite differences
 - Operators


Pattern Matching/Detection

- The optimal (highest) response from a filter is the autocorrelation evaluated at position zero


$$\max_{\bar{x}} C_{ff}(\bar{x}) = C_{ff}(0) = \int f(\bar{s})f(\bar{s})d\bar{s}$$

- A filter responds best when it matches a pattern that looks like itself
- Strategy
 - Detect objects in images by correlation with “matched” filter


Measure of correlation in time series at different lags




No lag, multiply and sum area




Measure of correlation in time series at different lags




Small lag, multiply and sum area




Measure of correlation in time series at different lags




Large lag, multiply and sum area



Measure of correlation in time series at different lags



lag, multiply and sum area




Reasoning


- Dot product of two normalized vectors a and b is maximum when both the vectors are equal.

$$a \cdot b = |a||b| \cos(\theta)$$


Match Filter Example



Trick: make sure
kernel sums to
zero




Match Filter Example




[Responses from match filtering]

Match Filter Example



[Peaks from match filtering]

Match Filter Example



Derivatives: Finite Differences

$$\frac{\partial f}{\partial x} \approx \frac{f(x+1, y) - f(x-1, y)}{2}$$

$$\frac{\partial f}{\partial y} \approx \frac{f(x, y+1) - f(x, y-1)}{2}$$


$$\frac{\partial f}{\partial x} \approx w_{dx} \circ f \quad w_{dx} = \begin{bmatrix} -\frac{1}{2} & 0 & \frac{1}{2} \end{bmatrix}$$

$$\frac{\partial f}{\partial y} \approx w_{dy} \circ f \quad w_{dy} = \begin{bmatrix} -\frac{1}{2} \\ 0 \\ \frac{1}{2} \end{bmatrix}$$


Derivative Example



$$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$



$$\begin{bmatrix} 0 & -1 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$



Other filters

- Prewitt

-1	0	+1
-1	0	+1
-1	0	+1

-1	-1	-1
0	0	0
+1	+1	+1

- Sobel

-1	0	+1
-2	0	+2
-1	0	+1

-1	-2	-1
0	0	0
+1	+2	+1

Convolution

- Discrete

$$g(x, y) = w(x, y) * f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

- Continuous

$$g(x, y) = w(x, y) * f(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} w(s, t) f(x - s, y - t) ds dt$$

- Same as cross correlation with kernel transposed around each axis

Convolution = Correlation

- The two operations (correlation and convolution) are the same if the kernel is symmetric about axes

$$g = w \circ f = w^* * f \quad w^* \text{ reflection of } w$$

Convolution: Properties

- Shift invariant, linear
- Commutative

$$f * g = g * f$$

- Associative

$$f * (g * h) = (f * g) * h$$

- Others (discussed later):
 - Derivatives, convolution theorem, spectrum...

Associativity

Example: Correlation

$$\begin{array}{c} [-1 \ 3 \ -3 \ 1] \\ \underbrace{\hspace{1.5cm}} \\ [1 \ -2 \ 1] \\ \underbrace{\hspace{1.5cm}} \\ [-1 \ 1] \ [1 \ -1] \ [-1 \ 1] \\ \underbrace{\hspace{2.5cm}} \\ [1 \ -2 \ 1] \\ \underbrace{\hspace{2.5cm}} \\ [1 \ -3 \ 3 \ -1] \end{array}$$

Associativity

Example: Convolution

$$\begin{array}{c} [1 \ -3 \ 3 \ -1] \\ \underbrace{\hspace{1.5cm}} \\ [-1 \ 2 \ -1] \\ \underbrace{\hspace{1.5cm}} \\ [-1 \ 1] \ [1 \ -1] \ [-1 \ 1] \\ \underbrace{\hspace{2.5cm}} \\ [-1 \ 2 \ -1] \\ \underbrace{\hspace{1.5cm}} \\ [1 \ -3 \ 3 \ -1] \end{array}$$

Computing Convolution

- Compute time

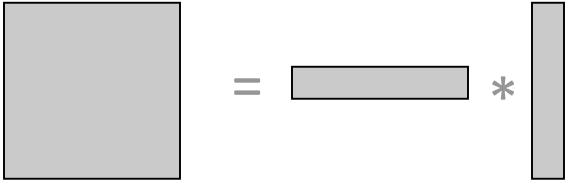
- $M \times M$ mask
- $N \times N$ image

$O(M^2N^2)$

“for” loops are nested 4 deep

- Special case: separable

Two 1D kernels

$$w = w_x * w_y$$

$$w * f = \underbrace{(w_x * w_y) * f}_{O(M^2N^2)} = w_x * \underbrace{(w_y * f)}_{O(MN^2)}$$

Separable Kernels


- Examples

- Box/rectangle
- Bilinear interpolation
- Combinations of partial derivatives
 - $d^2f/dxdy$
- Gaussian
 - Only filter that is both circularly symmetric and separable

- Counter examples

- Disk

- Cone



- Pyramid

Examples of Separable Kernels

Smoothing Filter

$$\frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} * \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \end{bmatrix} = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Sobel Filter

$$\mathbf{G}_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} * A = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} +1 & 0 & -1 \end{bmatrix} * A$$

Examples of Separable filter

Prewitt Filter

$$\begin{pmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{pmatrix} = ?$$

$$\begin{pmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{pmatrix} = ?$$

2nd order derivatives

Second derivatives:

$$\frac{\partial f}{\partial x} = [-1 \ 1] * f = f(x) - f(x - 1)$$

$$\begin{aligned} \frac{\partial}{\partial x} \frac{\partial f}{\partial x} &= [-1 \ 1] * [-1 \ 1] * f = [1 \ -2 \ 1] * f \\ &= f(x + 1) - 2f(x) + f(x - 1) \end{aligned}$$

$$\frac{\partial f}{\partial y} = \begin{bmatrix} -1 \\ 1 \end{bmatrix} * f = f(y) - f(y - 1)$$

$$\begin{aligned} \frac{\partial}{\partial x} \frac{\partial f}{\partial x} &= \begin{bmatrix} -1 \\ 1 \end{bmatrix} * \begin{bmatrix} -1 \\ 1 \end{bmatrix} * f = \begin{bmatrix} 1 \\ -2 \\ 1 \end{bmatrix} * f \\ &= f(x + 1) - 2f(x) + f(x - 1) \end{aligned}$$

Laplacian

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

$$= f(x+1, y) + f(x-1, y) - 2f(x, y)$$

$$+ f(x, y+1) + f(x, y-1) - 2f(x, y)$$




$$= f(x+1, y) + f(x-1, y) + f(x, y+1) + f(x, y-1) - 4f(x, y)$$

0	+1	0
+1	-4	+1
0	+1	0

Laplacian of a Gaussian

- We take an image and blur it a little using Gaussian function.
- Calculate the second order derivatives or the Laplacian.
- locates edges and corners that are good for detecting keypoints.
- Computation of the second order derivative is also extremely sensitive to noise, and the blurring helps.

Laplacian of Gaussian



0	+1	0
+1	-4	+1
0	+1	0


$$* \frac{1}{2\pi\sigma} \exp\left[-\frac{(x^2 + y^2)}{2\sigma^2}\right]$$

Nonlinear Methods For Filtering

- Median filtering
- Bilateral filtering
- Neighborhood statistics and nonlocal filtering

Median Filtering

- For each neighborhood in image
 - Sliding window
 - Usually odd size (symmetric) 5x5, 7x7, ...
- Sort the greyscale values
- Set the center pixel to the median
- Important:
 - Separate input and output buffers
 - All statistics on the original image



Median Filter

- Issues

- Boundaries

- Compute on pixels that fall within window

- Computational efficiency

- What is the best algorithm?

- Properties

- Removes outliers (replacement noise – salt and pepper)

- Window size controls size of structures


- Preserves straight edges, but rounds corners and features

Median vs Gaussian


Original




+ Gaussian Noise



3x3 Median




3x3 Box



Median Filtering

- Image model: piecewise constant (flat)



Order Statistics

- Median is special case of order-statistics filters
- Instead of weights based on neighborhoods, weights are based on ordering of data

Neighborhood
 X_1, X_2, \dots, X_N

Ordering
 $X_{(1)} \leq X_{(2)} \leq \dots \leq X_{(N)}$

Filter $F(X_1, X_2, \dots, X_N) = \alpha_1 X_{(1)} + \alpha_2 X_{(2)} + \dots + \alpha_N X_{(N)}$

Neighborhood average (box)

$$\alpha_i = 1/N$$

Median filter

$$\alpha_i = \begin{cases} 1 & i = (N + 1)/2 \\ 0 & \text{otherwise} \end{cases}$$

Trimmed average (outlier removal)

$$\alpha_i = \begin{cases} 1/M & (N - M + 1)/2 \leq i \leq (N + M + 1)/2 \\ 0 & \text{otherwise} \end{cases}$$

Piecewise Flat Image Models

- Image piecewise flat -> average only within similar regions
- Problem: don't know region boundaries



Piecewise-Flat Image Models


- Assign probabilities to other pixels in the image belonging to the same region
- Two considerations
 - Distance: far away pixels are less likely to be same region
 - Intensity: pixels with different intensities are less likely to be same region

Piecewise-Flat Images and Pixel Averaging

Distance (kernel/pdf)

$$G(\mathbf{x}_i - \mathbf{x}_j)$$


Prob
pixel
belongs
to same
region as
i



Distance (pdf)

$$H(f_i - f_j)$$

Prob
pixel
belongs
to same
region as
i



Bilateral Filter

- Neighborhood – sliding window
- Weight contribution of neighbors according to:

$$f_i \leftarrow k_i^{-1} \sum_{j \in N} f_j G(\mathbf{x}_i - \mathbf{x}_j) H(f_i - f_j)$$

$$k_i = \sum_{j \in N} G(\mathbf{x}_i - \mathbf{x}_j) H(f_i - f_j)$$

- G is a Gaussian (or lowpass), as is H, N is neighborhood,
 - Often use $G(r_{ij})$ where r_{ij} is distance between pixels
 - Update must be normalized for the samples used in this (particular) summation
- Spatial Gaussian with extra weighting for intensity
 - Weighted average in neighborhood with downgrading of intensity outliers

Bilateral Filtering




Gaussian Blurring




Bilateral


Results



Original




Noisy




Filtered


Checkerboard With Noise



Original




Noisy




Filtered

MRI Head







MRI Head




Fingerprint




Fingerprint




Results



Original



Noisy




Filtered

Results



Original




Noisy




Filtered


Results



Original




Noisy




Filtered


Fractal



Original





Noisy



Filtered

Piecewise Constant

- Several 10s of Iterations
- Tends to obliterate rare events



Texture, Structure

